SMALL PREPARED TEST SET / 26 OCT 2026

# A small rule-based test, measured honestly

The evaluation checks whether the software returns the outcomes prepared in advance for these examples. It does not estimate real-world accuracy.

## Test setup and results

The test compares 80 fictional equipment records with 3 real CISA notices (240 software-only comparisons). Fourteen asset/notice examples were labeled ahead of time: 4 likely affected, 4 outside the notice conditions, and 6 that need a human check.

| Measure | Result | Scope |
|---|---|---|
| Prepared examples with expected result | 14 of 14 | Hand-labeled synthetic examples |
| Affected examples correctly flagged | 4 of 4 | Only four in this small test |
| Missed affected examples | 0 of 4 | In this fixture only |
| Uncertain examples sent to review | 6/6 | Missing/old/unclear records |
| Results with source ID, link, and excerpt | 14 of 14 | All labeled examples |
| Time for software-only comparison | 1.01 ms total | One pass over 240 pairs |

How to read this: an affected example is a case labeled in advance as matching the notice. "Sent to review" means the software lacked enough support to give a firm result, so a person must check. A missed affected example is a false negative. These counts describe only the prepared test cases.

**Interpretation: these results show that OTProof passed this small, prepared test set. Four affected examples are far too few to predict how it will perform on a real operator's equipment list.**

## AI suggestions

No local AI model was run. Its ability to extract details or summarize notices has not been measured. The public-notice importer has software tests, but no diverse set of notices has been labeled by experts for comparison.

EXAMPLE REGISTER

# The test cases and what OTProof returned

The examples cover exact and outside version matches, spelling differences, a similar but different product, missing details, unknown model, old inventory, and duplicate assets. The table compares the expected result with OTProof's observed result.

| Asset | Advisory | Expected | Observed | Source linked |
|---|---|---|---|---|
| CV-HMI-001 | ICSA-24-226-06 | applicable | applicable | Yes |
| CV-HMI-002 | ICSA-24-226-06 | not applicable | not applicable | Yes |
| CV-HMI-003 | ICSA-24-226-06 | applicable | applicable | Yes |
| CV-HMI-004 | ICSA-24-226-06 | not applicable | not applicable | Yes |
| CV-HMI-005 | ICSA-24-226-06 | requires manual verification | requires manual verification | Yes |
| CV-HMI-006 | ICSA-24-226-06 | requires manual verification | requires manual verification | Yes |
| CV-HMI-007 | ICSA-24-226-06 | requires manual verification | requires manual verification | Yes |
| CV-HMI-008 | ICSA-24-226-06 | requires manual verification | requires manual verification | Yes |
| CV-HMI-009 | ICSA-24-226-06 | requires manual verification | requires manual verification | Yes |
| CV-OPT-010 | ICSA-24-256-24 | applicable | applicable | Yes |
| CV-OPT-011 | ICSA-24-256-24 | not applicable | not applicable | Yes |
| CV-APP-012 | ICSA-24-284-15 | applicable | applicable | Yes |
| CV-APP-013 | ICSA-24-284-15 | not applicable | not applicable | Yes |
| CV-OPT-014 | ICSA-24-256-24 | requires manual verification | requires manual verification | Yes |

## Expected results compared with actual results

| Expected \ Actual | Likely affected | Outside notice | Human review |
|---|---|---|---|
| applicable | 4 | 0 | 0 |
| not applicable | 0 | 4 | 0 |
| requires manual verification | 0 | 0 | 6 |

**The software does not guess based on similar-sounding product names. Human review is the intended result when equipment or notice details are unclear.**

# Security and availability checks

The test suite checks local software behavior and known edge cases. It does not certify the system or validate real OT outcomes.

## Checks included

Tests include security instructions embedded in advisory text; the AI helper is told to treat them as untrusted, and unsupported extracted details or quotes are rejected. Duplicate equipment IDs and macro-enabled spreadsheets are rejected. When a source refresh fails, OTProof can use a saved older copy and mark it stale. Permission tests check that users cannot approve or verify work outside their role. The workflow test records a human decision, approved plan, simulated evidence, and a separate evidence check.

## Performance scope

1.008 ms elapsed for one in-process pass over 240 pairs (0.0042 ms/pair on this developer machine). The measurement excludes database queries, HTTP ingestion, browser rendering, concurrency, cold starts, and human review. Operator review time was not instrumented or measured.

## Limits

Small synthetic sample; three Rockwell CISA fixtures; no blinded ground truth; no live local model evaluation; no industrial vendor/model breadth study; no penetration test; no production load test; and no human time study. The test client emits a Starlette warning that the installed HTTPX integration is deprecated, though tests pass. Docker Compose could not be run because a daemon was unavailable.

## Next evaluation protocol

Recruit engineers from several operators. Ask them to label source-reviewed examples across vendors and versions before they see OTProof's results. Compare disagreements and report how often the system finds a relevant affected product, how often it misses one, how often it asks for human review, whether each result links to its source, and minutes spent per review. Test the rule-based matcher first. Evaluate AI extraction separately by comparing its suggested fields with the source and human labels. Also test unavailable sources, large spreadsheets, hostile text, permissions, and deployment security.

**The current result is a build check, not a customer study, security certification, safety case, or verified reduction in operational risk.**